

CMSC 216

INTRODUCTION TO COMPUTER SYSTEMS



AKILESH PRAVEEN
DR. ILCHUL YOON • SPRING 2020 • UNIVERSITY OF MARYLAND
<http://cs.umd.edu/class/spring2020/cmssc216/>

Last Revision: March 16, 2020

Table of Contents

1	Notes & Preface	3
2	Week 1 - Introduction to CMSC216	3
	Overview	3
	Grace	3
	Useful UNIX Commands	4
	Machine	4
3	Week 2	5
	The Math Library	5
	Using Emacs	5
	Debugging	6
	GDB	6
4	Week 3	6
	Comma is an Operator	6
	Identifier Scope	7
	C Program Memory Organization	7
	Storage	8
	Linkage	8
	Enumerated Types	9
	Implicit Type Conversion and Casting in C	9
5	Week 4	10
	Pointers & Memory Maps	10
	Example - Integer & Integer Pointer	11
	Example - Multiple Pointer Types	11
	Example - Pointer To a String (Char Array)	12
	Lab Examples	13
	Example from Lab - ptr_review.c	13
	Example from Lab - ptr_add_sub_overview.c	14
	Example from Lab - str_review.c	15
	Using getchar() and putchar()	15

6	Week 5	16
	Grep - A 'CTRL-F' From the Command Line	16
	Why is it called that?	16
	Why it's useful	16
	Memcpy, Memmove, and Memset	17
7	Closing Thoughts	17

1 Notes & Preface

This is a compilation of my notes for CMSC216 as a TA for the Spring 2020 offering of the course at the University of Maryland. All content covered in these notes was created by Dr. Ilchul Yoon and Dr. A.U. Shankar at the University of Maryland.

The actual content of this note repository is the content that I cover as a TA during my discussion section, combined with my personal insights for the course. I took this course with Nelson Padua-Perez in the Spring 2019 offering, so some of the notes that I'll drop in here are from my own notes when I took the course in 2019. As such, I would like to attribute certain code examples, analogies, and more to Mr. Perez. I believe that together, these will serve as great **supplementary material** for CMSC216, but I would still highly recommend attending all of your lecture and discussion sections to achieve success in CMSC216.

The notes template in use is Alex Reustle's template, which can be found on his github at the following location: <https://github.com/Areustle/CMSC351SP2016FLN>

I maintain this repository and as such, take responsibility for any mistakes. Please send errors to apraveen@cs.umd.edu

2 Week 1 - Introduction to CMSC216

CMSC216 is where you learn how a computer works on a much lower level than you've experienced before. There are 3 main components that the course will explore.

Overview

- **UNIX** Threads, processes, and pipes as the building blocks of much bigger applications. We will be working with the UNIX operating system on the development environment at grace.umd.edu
- **C** is a high-performance language that works at a much lower level than Java. Things like memory management and advanced data structures are left up to the user. We'll cover concepts like memory management, pointers, and system calls.
- **Assembly** is even lower-level than C, and studying it will reveal how processors process instructions, store data, and maintain a stack and a heap. It's the lowest level you'll go in this class. For this semester's 216, you will be using MIPS assembly.

Grace

In this class, we will be using the Grace system to do all of our work. It's a little confusing to understand at first, so here's my way of thinking about it. In CMSC132, we did all of our work on our own computers. We pulled the skeleton code for the projects from the 132 website/repository, edited the code on our computers, and then uploaded our code to the submit server (via Eclipse) in order to test it.

In CMSC216, we have been given access to this big computer that UMD CS owns known as Grace. You, as a student, have been given a small chunk of that machine to call your own (for the semester). In this class, we will access your files on the Grace system using a program known as `ssh` (that's how MobaXTerm works) and do all of our editing + running code on Grace itself. In fact, we will also be submitting our projects from Grace to the UMD CS submit server.

Here are the relevant links for getting it all set up. You'll need to setup Grace and `gcc` (the C compiler that we'll be using within Grace).

- <http://www.cs.umd.edu/~nelson/classes/resources/GraceSystem.shtml>
- http://www.cs.umd.edu/~nelson/classes/resources/setting_gcc_alias.shtml

Useful UNIX Commands

Although the UNIX environment may seem confusing at first, learning it is essential to navigating the Grace environment. Below are some of the basic commands that you may find useful when getting started.

- **ssh** → If you are not using MobaXTerm, you will have to access grace using the **ssh** command. For the purpose of logging in for CMSC216, I recommend adding the **-y** flag in order to bypass the warning it will give you. E.g. **ssh -y yourdirectoryid@grace.umd.edu**
- **ls** → The **ls** command lists all the files in your current directory. You can use the **-l** flag to get more detailed information. E.g. **ls, ls -l**
- **cd** → The **cd** command changes the directory you're currently in, mainly to directories that you can see with **ls**. Typing **cd ..** will navigate one directory 'up' from your current directory, and **cd** without anything else will return you to your home directory. E.g. **cd 216public**
- **pwd** → This command displays your current directory. Useful for finding out where exactly you are in the UNIX file hierarchy. E.g. **pwd**
- **cp** → Copies files. If you use the **-r** flag, you're telling the command to recursively copy. If you want to use **cp** on directories, remember to use that flag.
- **rm** → This command stands for 'remove'. It can be used to remove singular files, or can alternatively be used with the **-r** flag to recursively remove directories. E.g. **rm hello.c, rm -r project1** (project1 would be a folder).
- **., .., ., and /** → These abbreviations are pretty important. They can be used to navigate a filesystem in Unix and generate some clever commands. In order, they mean 'current directory', 'parent directory', 'user home folder', and 'root directory'. Below are some examples.
 - **cp *.c ../** → Copies all files that end with **.c** to the parent directory.
 - **cd /** → Changes directories to the root directory.
 - **cp -r /216public/projects/project1 .** → Recursively copies (this means that it copies directories as well as files) the project1 directory and everything in it into the current directory.

Lots of these UNIX commands are super useful once you get to know them, but it may be hard becoming acquainted with how they work from the outset. It's a far cry from the GUI you had in CMSC132, so here are a few tips.

- If you're just starting out and still need a graphical representation of the filesystem, I'd highly recommend setting up **MobaXTerm**. The program provides just a little more graphical representation than just a pure terminal, and allows you to navigate the Grace filesystem more freely. I like to think of it as training wheels as you get acquainted with Grace.
- I'd highly recommend getting used to making folders, deleting folders, deleting files, and navigating up and down through the filesystem with rapid sequences of **ls** and **cd**. As with all things, practice makes perfect, and pretty soon you'll be a command line wizard.

Machine

A computer is composed of several parts, but a great way to think about it is a few main components connected by a **bus**.

- **Memory** can just be thought of as a contiguous array of bytes. At the end of the day, this is the stuff that has to be written to/read from.
- **I/O Devices** are connected to the CPU via a bus, like mentioned above. By performing read/write operations to the right adaptor, the CPU is able to interface with different I/O devices.
- **CPU** is the central processing unit of the computer. It handles computational operations (arithmetic, logic, etc.) and interfaces with the memory and I/O devices via the bus. The CPU is also responsible for performing the **fetch-execute cycle**.

A bus is like one main connector that's responsible for making sure the CPU, memory, and I/O devices are all able to interface with each other.

Note that in this course, we won't be going too in-depth into hardware (that's more Computer Engineering), but it's great background knowledge to have as you approach this class, which is why I have included it here.

3 Week 2

The Math Library

We won't be using the math library much in C, but for the times that we do, just remember this one simple flag that we add to the gcc command. As an example, if you try to write some code that includes the math library like below, you'll find that it won't compile with a regular gcc command.

```
1 #include <stdio.h>
2 #include <math.h>
3
4 int main() {
5     double value;
6
7     printf("Enter a number: ");
8     scanf("%lf", &value);    /* Notice the use of %lf */
9
10    printf("sqrt %f: \n", sqrt(value));
11    printf("power of 2: %f\n", pow(value, 2));
12    printf("sin: %f\n", sin(value));
13
14    return 0;
15 }
```

Remember that the `-lm` flag essentially enables us to use the math library. In other words, if you want to compile the above file and have it work properly, (let's assume it's called `math_example.c`) then you'll want to compile it using the following command.

```
gcc -lm math_example.c
```

Using Emacs

Most of the instruction for this course will be done in `emacs`, a highly versatile text editor that you can use in GUI form or from the command line. It's always an option to use other text editors in this class, but I would recommend using `emacs`, as it's what all the in-class demos are in. There is a way to setup IDEs like Visual Studio Code to function with Grace, but I won't cover them here. I believe that although graphical IDEs have their advantages, you'll get plenty of experience with them in CMSC330 and CMSC4XX, so for now, develop your skills in a command line editor like `emacs` or `vi`.

For your benefit, here are some basic commands in `emacs` that I've found useful over the time that I took 216.

Note: When I indicate to type M, that means you need to press the 'meta' key. On most machines, the 'meta' key is the 'alt/option'. When I indicate to type C, I mean the 'control' key. The reason I'm using this notation is because it's the same notation that online guides use to describe `emacs` shortcuts.

- **C-x C-s** → Saves the file you're working on. Remember to do this frequently on Grace, as you can't guarantee that your connection to Grace will stay intact.
- **C-x C-c** → Closes the file that you're working on. If you haven't saved, it will prompt you to save.
- **C-x u** → Undo the previous command that you ran.
- **C-s** → Search forwards (this will search for text that'll be ahead of where your cursor is now.)

- **C-r** → Search backwards (this will search for text that'll be behind where your cursor is now.)
- **C-l** → This command will center the window around your cursor. A great technique when you have large C files that you're editing.
- **M-x column-number-mode** → Shows column numbers. Useful if you want to check if you're above the 80 character limit.

Debugging

There are three main debugging tools that we use in 216: Valgrind, GDB, and splint. For now, we won't focus too much on Valgrind, as it's more oriented towards helping programmers get rid of memory leaks and other memory-related issues. We will focus on GDB and Splint.

GDB

GDB Is the C equivalent of the Eclipse Debugger. It lets you do everything that the Eclipse Debugger allowed you to do in CMSC131 and CMSC132. The only real drawback here is that it's all done from the command line, so the graphic part of the interface is a little lacking. However, it's an essential tool that I'd highly recommend using to figure out errors in your code.

Online references will tell you that there are a lot of commands that you need to know to effectively use GDB, but here are some of the ones that I've found useful.

- **q** → exits gdb. Useful.
- **start** → starts running your code with a temporary breakpoint at the first line of `main()`. This allows you to set more breakpoints before the code actually starts executing.
- **l** → lists the code that you have.
- **b** → typing `p` with a number next to it sets a breakpoint at a line. E.g. `b 3`
- **n** → the equivalent of step over in the Eclipse debugger
- **s** → the equivalent of step into in the Eclipse debugger
- **c** → will continue running your code until the next breakpoint
- **p** → will print the value of an expression or a variable. E.g. `p valid_character('x')`.

In order to start GDB, you'll first need to compile your C code into an `a.out` file. Not only that, but I would recommend that you compile your code with the `-ggdb` flag, to ensure that GDB initializes your program correctly. In order to run GDB with your newly compiled program, remember to just type `gdb a.out`

4 Week 3

This week we go over a lot of general C-specific programming concepts in discussion, and that material is heavier than what we usually do in discussion. In that sense, I'll try and go over the more basic stuff that I think will be highly useful as you work on your projects.

Comma is an Operator

The comma in C is an operator. The best way to think about this in use is when you're declaring multiple variables at once, like when you say `int i, j = 2`.

Remember, commas are **also** used as separators in C. A great example would be if you're giving a function multiple parameters, like in `printf("%d and %d", i, j)`. When you consider the comma as an operator in C, it's always important to understand where it's an operator vs. where it's a separator.

Although we don't think about the comma operator quite a lot, one of the main reasons for understanding it would be initialization of multiple variables in a loop. Take a look at the following example from my notes (from a previous offering of the CMSC216 course).

```

1 // Comma Operator Example by Nelson Padua-Perez
2
3 for (j=0, k=10; j<=limit; j++, k+= 10) {
4     printf("j->%d, k->%d\n", j, k);
5 }

```

Notice how you initialize and increment multiple variables within a single for-loop.


Identifier Scope

Scope exists in C in a similar way that it does in other languages. All you have to remember is that if you declare variables within code blocks, they won't be accessible outside those blocks. In that regard, this phenomenon is quite similar to how Java handles scopes.

C Program Memory Organization

As we delve deeper into systems-level programming, it's important to visualize how C actually manages the memory that your programs use. The interesting part about this is that this diagram is an exact representation of system memory, so you're finally able to see 'under the hood' of your programs.

You can see that the lowest address is represented by `0x0` and the highest address is represented by `0xffffffff`. These addresses are actual locations in memory, represented in hexadecimal format (hence `ffffff` being the highest address in the representation).



Your program is allocated a certain block of memory- within it are the following 4 components. Keep in mind that this too, is an abstraction. You can further explore how programs are represented in memory in classes like CMSC411, but this is just about as far as we'll go in 216.

- **Text** is where the code for your program goes. It's really not much more complex than that.
- **Data** is where global variables and variables that are static belong.
- **Heap** is where dynamically allocated memory lives. In Java, this stuff was managed for you. In C, you will have to manage it yourself, allocating memory and effectively increasing the size of the heap if you need more space while your program is running, and deallocating (freeing) memory to decrease the size of the heap. More on this when we discuss dynamically allocated memory.
- **Stack** is where local variables and function parameters live. It grows downwards (eventually meeting the heap and causing a stackoverflow) as functions are called. If you'll think back to 'stack frames' from recursion in CMSC132, this is the exact same concept.

Storage

There are two types of ways variables are stored in C- automatic and static. This basically goes hand-in-hand with block scopes and file scopes, but the important takeaways are these. First of all, in the example below, after the function `foo` is called, the variable `n` is thrown away.

```
1 int foo(int k) {  
2     int n = 216;  
3     return n;  
4 }
```

In that regard, the variable `n` has automatic storage. If a variable has static storage, it basically exists throughout the duration of your program's running time. Such variables are initialized only once.

An important note: `static` in C does not mean the same thing as it does in Java. Here are the two main things that I think are worth remembering about static variables in C:

- Static variables need not necessarily be initialized. If you don't bother initializing a static variable (you still have to declare it- this is not Python, language of the heathen) it will automatically initialize to zero.
- Static variables retain their values between function invocations. In other words, they are not stored using automatic storage.

```
1 // example from Nelson Padua-Perez  
2  
3 void compute_static(int x) {  
4     static int value = 100; /* What would happen if we don't initialize ↵  
5        it? */  
6     printf("(static) x: %d, value: %d, sum: %d \n", x, value, value + x)↵  
7     ;  
8     ++value;  
9 }
```

In the example above, if you called `compute_static` twice, then your output would be (static) x: 1, value: 100, sum: 101 and (static) x: 1, value: 101, sum: 102, as 'value' would retain its data between function calls.

Linkage

Linkage is essentially the science behind having C code spread across multiple files and making sure it all compiles and works properly.

We want to sometimes split code between multiple files for organizational purposes. Currently, the projects you're working on are small, but in order to make your programs versatile, modular and better organized, it's a great idea to split code between files.

When you attempt this, there may be issues that follow. For example, you may encounter a situation where you want to name a function `print_sum()` in two files. How would we deal with such a duplicate?

Problems of this sort can be solved by adjusting the **linkage** of these functions.

For actual code examples, please check the linkage-examples in the 216public directory. They're extremely thorough. My goal here is to provide a quick few tips on what I think are the most important parts.

Essentially, there are three types of linkage that you should remember to guide you through writing code in multiple C files.

- **None** → No linkage. This is how you usually declare your variables, and as you'd expect, doesn't do anything special in regards to linkage. Think of it this way: A variable with no linkage belongs to a single function, and cannot be shared. In other words, there is *only one copy per declaration*.
- **Internal** → Internal linkage is just a fancy way of saying you're using the **static** keyword. All declarations of a single identifier in file refer to the same thing. In other words, there is *only one copy per file*.
- **External** → External linkage is signified by the **extern** identifier, and it basically means that a name can only refer to a single entity in your entire program. In other words, there is *only one copy per program*.

Enumerated Types

Enumerated Types, or enums, in C are pretty useful, and quite comparable to their equivalents in Java. The best way to understand enums (in my opinion) is to think of examples. Some good ones are an enum for the days of the week (Monday, Tuesday, etc.), seasons (Summer, Spring, Fall, Winter), or even suits in a deck of cards (Spades, Clubs, Hearts, Diamonds). Below is an example of the latter.

```
1 // example from Nelson Padua-Perez
2 #include <stdio.h>
3 int main() {
4     enum Suit { SPADES, HEARTS, DIAMONDS = 42, CLUBS};
5     enum Suit suit1, suit2;
6     suit1 = SPADES;
7     suit2 = CLUBS;
8     if (suit1 < suit2) printf("Spades are first.\n");
9     else printf("Clubs are first.\n");
10    printf("Spades = %d, Clubs = %d\n", suit1, suit2);
11    return 0;
12 }
```

The functionality here is pretty basic, but one thing that I think is worth remembering (and quite nifty if you can use it well) is that enum representations are based in integers. This means that, for example, you can get away with adding the month enum for January (0) and the month enum for February (1) and end up with February (1).

Again, the code above is a great example of how you can leverage the integer-like characteristics of enums.

Implicit Type Conversion and Casting in C

Switching between data types is pretty similar to how it was in Java, but here's a quick review of the stuff that matters. As you write your projects, you'll realize these things, but it's important to remember when it's a good idea to cast and when it isn't. Here are some general tips for you.

- ints are the smallest data type. If you need to describe something as a number and want to be as efficient as possible without getting fancy, the good old integer is the best way to go.
- We can also cast in C, and it works almost the same as it did in Java. Just remember, in Java we had the concept of wrapper classes that allowed us to do fancy things with certain data types. In C, we don't enjoy that luxury, so we are restricted to just basic data type casting. Below is an example.

```
1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 int main() {
```

```

6     float x = 2.98;
7     int y = (int)x;
8 }

```

- That works exactly as you think it does. It converts 2.98 to 2 as it would in Java. Remember, don't overthink it, and don't try to call any wrapper class methods that you remember from Java. As long as you keep that in mind, you should be good to convert between data types in C.

```

1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 int main() {
6     int x = 2000000000;
7     long result_long;
8
9     printf("Value of x: %d\n", x);
10    printf("Multiplying by 3 (with %d format): %d\n", 2000000000 * 3);
11    printf("Multiplying by 3 (with %ld format): %ld\n", 2000000000 * 3)↵
12    ;
13    printf("Multiplying by 3L (with %d format): %d\n", 2000000000 * 3L)↵
14    ;
15    printf("Multiplying by 3L (with %ld format): %ld\n", 2000000000 * 3↵
16    L);
17
18    result_long = 2000000000 * 3; /* Does it solve the problem? */
19    printf("Storing result in long type variable: %ld\n", result_long);
20
21    return 0;
22 }

```

The above example from Nelson isn't that basic, but I feel like it gives you a good insight into how type conversion can find use. Give that example a try to see a cool application of using multiple data types to handle larger values.

5 Week 4

This week we cover pointers, a few functions in C that you may find useful, and GDB in emacs. The main focus of these notes will be pointers, and chances are that you've seen a lot of this stuff in lecture as well. Make sure to take some time to try out the examples that we've got for you so you understand the basics of how pointers work, as they're a fundamental part of C.

Pointers & Memory Maps

Let's go over pointers in C. You may have already covered this subject in lecture, but I'd like to point out some of the nuances that helped me understand pointers when I was taking 216.

First of all, take note that pointers are just another type of variable. Just like you have `ints` and `chars` in C, which take up a certain amount of space and store a certain type of data, a **pointer** is a data type that stores an **address**.


There are a bunch of ways to think of pointers, but I think the easiest way to understand them is to use memory maps. Think of them as a tool to help us better understand how pointers work- they are essentially just visual representations of memory in C.

I think that pointers and memory maps go hand in hand in 216, so I'll include some examples (some of my own, plus the examples we go over in discussion) that I think will help you become proficient with both pointers and memory maps.

As a side note, you can take a look at Nelson's sample memory map online if you need some extra guidance. (This should have been covered in discussion).

<http://www.cs.umd.edu/~nelson/classes/resources/MemoryMapExample.pdf>

Example - Integer & Integer Pointer



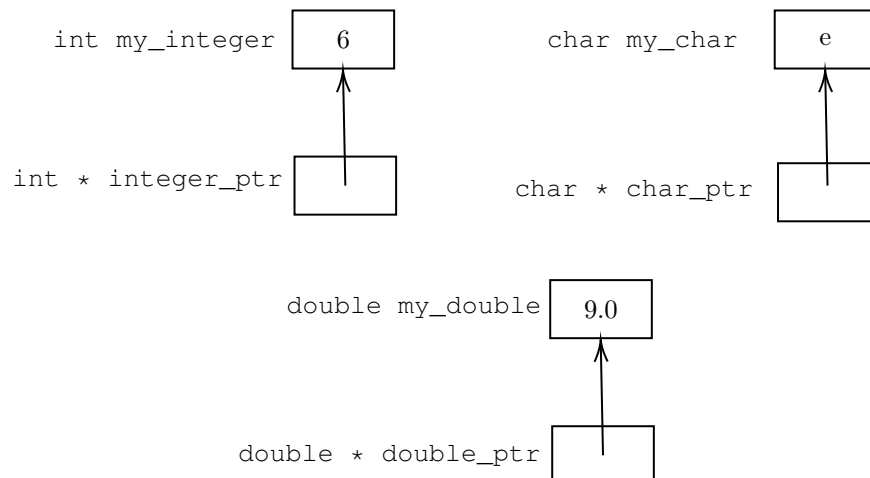
```
1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 int main() {
6
7     int a = 5;
8
9     int * b = &a;
10
11     return 1;
12 }
```

This is about as simple as we can get with pointers. There are a variety of types of pointers that exist (one for each data type in C), but just remember that they're essentially just variables that store addresses.

In this example, we can see that **a** is an integer, and **b** is an integer pointer. Although I've drawn an arrow from the inside of **b**'s box to **a**'s box, don't let that confuse you.

Think of it like this- **a contains** the integer value 5. **b contains** the address of **a**. By convention in C, we say that **b** points to **a**. We just show this by drawing an arrow that starts in **b**'s box and points to **a**.

Example - Multiple Pointer Types



```


1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 int main() {
6
7     int my\_integer = 6;
8     double my\_double = 9.0;
9     char my\_char = 'e';
10
11     int * int\_ptr = &my\_integer;
12     double * double\_ptr = &my\_double;
13     char * char\_ptr = &my\_char;
14
15 }

```

Here's a similar case to up above, but I just wanted to demonstrate that there are different types of pointers. Now, keep in mind that all of these pointers essentially hold addresses, and it's not like the address of a double looks much different from the address of a character or the address of an integer.

If you're wondering why C is so specific and asks you to define the type of pointer, the answer lies in how we will treat the data that's within the pointer. Sure, it may be that all pointers hold addresses, but what happens if we try to add the contents of `double_ptr` and `integer_ptr`? If C only had one pointer type and we tried to add the contents of those two pointers together, there would be no way of knowing that we made a mistake until runtime. In that sense, C maintains different types of pointers to ensure type compatibility. The same address could be given by the C memory manager to an integer pointer or a double pointer, but in order to make sure that you're treating whatever is stored at that address in a type-compatible way, C makes sure to note the type of what you're pointing to.

Example - Pointer To a String (Char Array)



```

1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 int main() {
6
7     char my_string[5];
8
9 }

```

Finally, here's a look at how we would store a string. I picked a string because it's essentially an array of characters, so we get to see how both are represented in memory maps.

Here, don't let the notation confuse you. Although I've declared the string `my_string` in special notation, it's still essentially a pointer to a character. In this case, `my_string` is a pointer to the first of 5 characters that C has allocated as NULL for us. I've taken the liberty to fill the allocated blocks in as null bytes.

Lab Examples

I'll also go over the examples that we went over in lab, but a little less in-depth, as they're usually a bunch of concepts put together. We'll focus on what I think are the important portions of each example.


Example from Lab - ptr_review.c

Here, we'll talk a little bit about ptr_review.c
(This file can be found at /216public/labs/Week4/lab1)

This is just going over the basics of pointers, and it has a few functions that demonstrate a few things, but I'd just like to go over a few of the questions posed in the actual file.

```
1 // example from Nelson Padua-Perez
2
3 int main(void) { /* notice use of void in main */
4     float *p, *m; /* have garbage value */
5     float pressure; /* has garbage value */
6     int area = 10;
7     int a[3] = {777, 888}; /* missing value? */
8
9     p = &pressure; /* & returns address */
10    m = p; /* both m and p point to the same entity */
11    printf("Value1 %.2f\n", *m); /* are we ever getting a segmentation ↔
12                                   fault?*/
13
14    ...
15    return 0;
16
17 }
```

- Using the keyword void in main essentially means that your program will be taking no arguments. That's the long and short of it.
- When we define p and m as pointers and don't assign anything to them, they essentially contain garbage values. If you want a visual representation of that, just imagine two pointer variables with arrows pointing into the unknown. We don't know what they're pointing to, nor do we want to find out.
- It's the same deal if we define a float without assigning it a value- it contains a garbage value.
- When they set m equal to p, they're making it so both pointers are pointing to the same variable. If that confuses you, think of it the other way- pointers contain addresses, and it just so happens that after executing m = p;, both m and p contain the same addresses.




- Finally, when it asks if we are ever getting a segfault, the short answer is **maybe**. In C, dereferencing a pointer that we have not yet initialized is considered **undefined behavior**. It could provide us with garbage data, give us a segfault because we tried to access corrupted data, or give us a segfault because we tried to access data locked off by the system. We don't really know what will happen in this case, so we're calling it undefined behavior. In grace, variables that aren't initialized are given a value of 0 or NULL, so we won't see this effect here. However, running in any other C environment will yield undefined behavior.

Example from Lab - ptr_add_sub_overview.c

Here, we'll talk a little bit about ptr_add_sub_overview.c
(This file can be found at /216public/labs/Week4/lab1)

This example is all about pointer arithmetic, and it relies on the fact that you understand that arrays are stored in contiguous memory. Let's think about the following example. If you had an array that was represented in C memory like this:



In this case, since arrays are stored in contiguous memory, so essentially what we are claiming with pointer arithmetic is that, if we dereference `p` now, we will get the number 1. If we **add** 1 to `p` (the actual pointer) and then dereference it, we will get the number 6. The file explores similar examples. Here are some highlights.

- Just like we discussed earlier, here's an application of simple pointer addition. As a reminder you can add numbers other than 1.

```
1 // example from Nelson Padua-Perez
2
3 char name[MAX] = "The House is Blue";
4 char *p = name, *q;
5 int i;
6
7 /* You can add and subtract integer values from pointers. */
8 /* For example, if you add one to a pointer to a character */
9 /* array, the pointer will now be referring to the next */
10 /* character. You can add any integer value (not just one) */
11
12 /* Printing the string using pointer arithmetic */
13 while (*p != '\0') {
14     printf("%c", *p);
15     p = p + 1;
16 }
17
18 }
```

- You can also take advantage of the fact that arrays are stored in contiguous memory by subtracting pointers to find 'distance' between them. Note that this only works with pointers of the same type.

```
1 // example from Nelson Padua-Perez
2
3 /* You can tell how many elements are between two pointers */
4 /* by subtracting pointers */
5 p = name + 1;
6 q = &name[5];
7 printf("Elements #1: %ld\n", q - p);
8 printf("Elements #2: %ld\n", p - q);
```

- Finally, you can leverage pointer arithmetic to help you index arrays as well. Here's an example of that below.

```

1 // example from Nelson Padua-Perez
2
3 /* Indexing is a pointer operation */
4 printf("Indexing as pointer operation\n");
5 p = name;
6 for (i = 0; i < strlen(name); i++) {
7     printf("%c\n", p[i]);
8 }

```

Example from Lab - str_review.c

Here, we'll talk a little bit about `str_review.c`

(This file can be found at `/216public/labs/Week4/lab1`)

This example is pretty light compared to the rest- and it is just a review of how strings are stored in C. The main overarching concept you need to understand here is two things:

- Strings are not given an actual data type in C. They are simply arrays of characters with a small caveat.
- That being said, strings are always stored in a certain way. They are a character array terminated with a null byte. (No null byte at the end means you don't have a string- you have a regular old character array)

Take a look at my String example above for the memory map representation.

Using `getchar()` and `putchar()`

The two functions `getchar` and `putchar` are pretty curious, in that we have much more functional replacements for them- `scanf` and `printf`, respectively. However, learning these is a cool way to prep yourself for how basic I/O in assembly works, so I think that it's worth it to at least gloss over these for now.

Let's look over the code provided for us in discussion and touch on the main points.

```

1 // example from Nelson Padua-Perez
2
3 #include <stdio.h>
4
5 #define MAX_LEN 80
6
7 int main() {
8     char value[MAX_LEN + 1];
9     int letter; /* Why integer? */
10
11     printf("Enter a letter: ");
12     scanf("%1s", value);
13     printf("Value entered: \"%s\"\n", value);
14     getchar(); /* getchar() reads a single character; why we need it? ↔
15                */
16     printf("Enter a letter: ");
17     letter = getchar();
18     printf("Letter entered: ");
19     putchar(letter); /* putchar() prints a single character */
20     printf("\n");
21
22     /* try ungetc to put characters back */
23
24     return 0;
25 }

```

- First of all, both `getchar` and `putchar` deal with integers, despite the fact that they are meant to take in/print characters. Don't let this confuse you, they're simply storing them by the ASCII value.
- Both of these get and print a single character, and in my opinion, there's no real reason to need them except in very special cases, but this is how I/O will be conducted in Assembly, so I think it's worth taking a look at this now.
- Your main takeaway from this should be that `getchar` and `putchar` are functions that we can use to do I/O in C, and even though they're a little more crude than we'd like for most applications, they still exist, and are helpful tools when we're trying to understand Assembly.

6 Week 5

Grep - A 'CTRL-F' From the Command Line

When working with the command line, we have the unique opportunity to see older versions of computer tools that we are accustomed to today. In modern environments, if you want to find something on a webpage, textbook, or even in your `.java` file in Eclipse, the first thing that probably comes to you head is the command 'CTRL + F'. In a command line environment, the command that preceded this functionality is known as `grep`.

Why is it called that?

The name of the command itself has an interesting origin. The most basic text editor on UNIX systems is regarded by many as `ed`, and on that text editor, one was able to globally search the file for a regular expression (which you'll learn more about in CMSC330), then print what was found using the command '`g/re/p`'. This gave way to the name "grep".

Why it's useful

As you'd imagine, `grep` can be used to simply search the files we have for keywords. Let's take a look at some examples. You can follow along if you head over to `216public/labs/Week5/lab 1/grep_example`. Let's take a look at the text files that we will be searching through, as examples.

```
1 The college is in
2 the east coast.
```

data.txt

```
1 The project is about hashing,
2 files, structures,
3 pointers
4 and dynamic memory allocation (and more pointers).
```

summary.txt

These two files are in the same directory, and for the purpose of the examples I'll go over, let's assume that we're currently in the directory that contains both these files.

`grep` works like this: you provide it a key phrase and a file location, and it'll take care of the rest. If you want more technical information on how `grep` commands should be structured, I encourage you to take a look at `man grep`.

If you execute the command `grep college data.txt`, then `grep` will print out the line that it found your keyword on. (the output for that command will be `The college is in.`)

Where `grep` really shines is when you want to mix in some of the cool UNIX keywords we've been learning. As a quick example, let's say you wanted to search for all the occurrences of 'is' in all the text files you had in the file. To do that, you'd simply execute the following command.

```
grep is *
```

That would yield the following:

```
data.txt:The college is in
summary.txt:The project is about hashing,
```

In a more practical example, let's think about how you could use this when writing your projects. Let's say you had a particularly tough project with 20 public tests. You're failing a bunch of them, but you suspect it's because the tests are calling a function you know you haven't implemented properly yet, named `get_classroom_number()`.

Assuming that public test files are named as they usually are in this class, and that you're in your project directory, if you wanted to figure out which public tests were testing for the `get_classroom_number()` function, all you have to do is cook up a `grep` command to do that for you. Here's what we'd be looking at in this case:

```
grep get_classroom_number() public*
```

This would search for the keyword `get_classroom_number()` in every file that started with 'public', which is exactly what we want. (Remember your UNIX special characters!). Additionally, here's one extra little trick that might make grepping a little bit easier- if you want to see the line numbers that your searches actually appear on, go ahead and use the `-n` flag when you run `grep`. The previous example would then look like this:

```
grep -n get_classroom_number() public*
```

Memcpy, Memmove, and Memset

7 Closing Thoughts

This document will be updated frequently as we progress through CMSC216. Please send errors to apraveen@cs.umd.edu